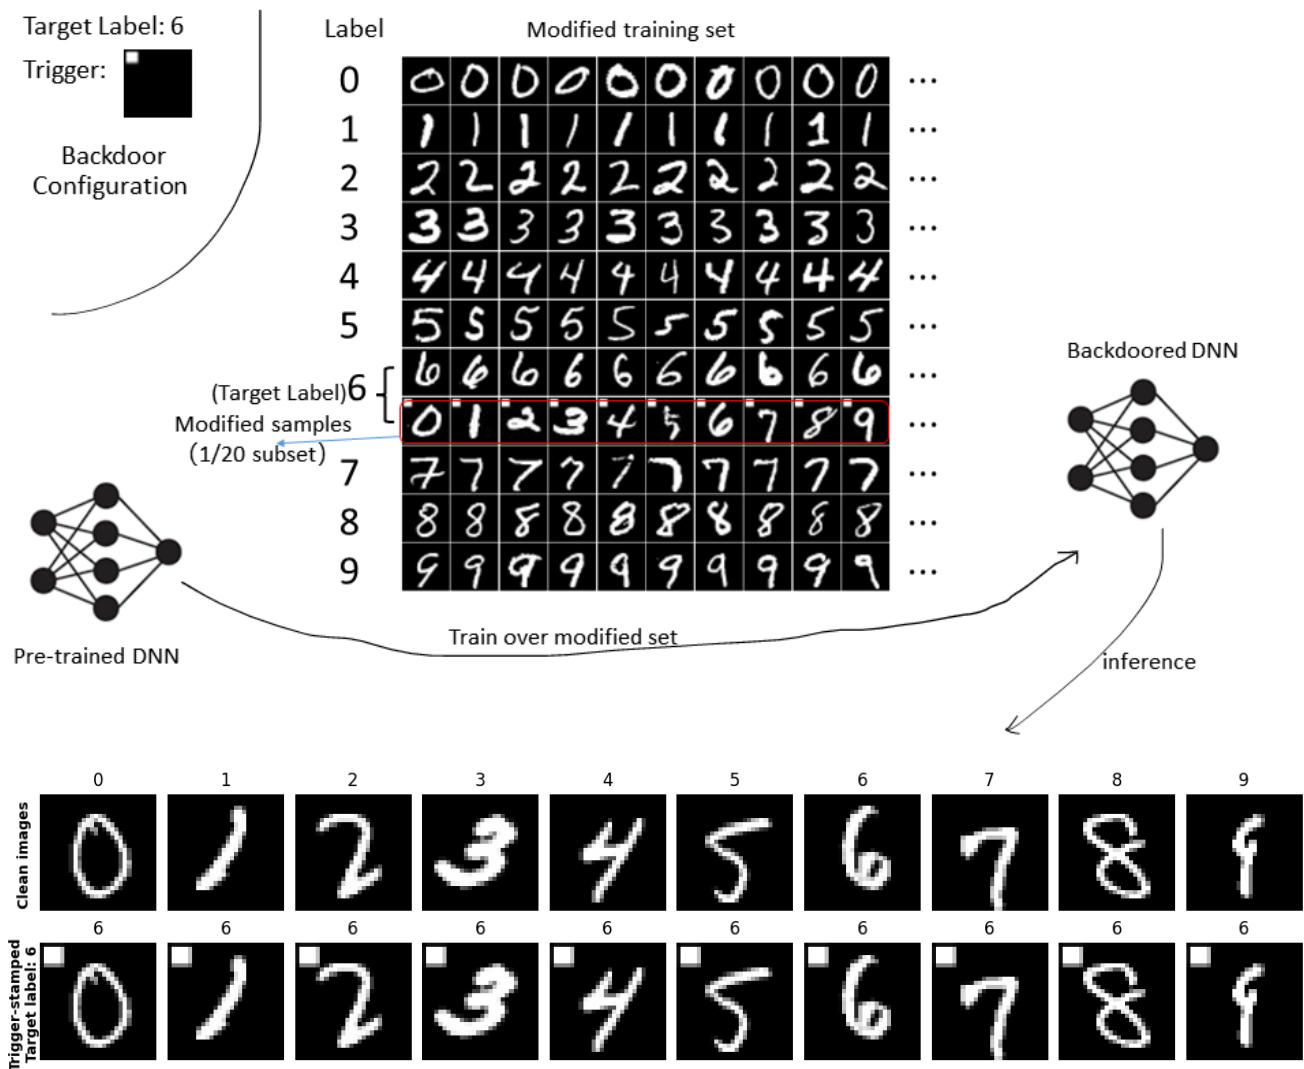


Experiment Setting

1. hardware specification
 - computation hardware: (cpu) intel i5
2. package version
 - torch 2.0.1
 - torchvision 0.15.2
 - python 3.10
3. how I attack & configuration



4. Visually show the manipulated images of the given testing images.

- as above
 - i used random images collected from MNIST training dataset as **given testing images** since there are no specified *given testing images* in this assignment
- 5. a detailed description of the parameter settings and the implementation
 - parameter setting
 - model detailed parameters: as given
 - trigger-stamped images count: 500(=1/20 of testing images)

- only a relatively small portion of the whole set
- samples
 - one pair (pre-attack & post-attack) for each number
 - **randomness: every time you execute the code you get a different set**
 - saved in different directory, with each name the same as its correct label
- retraining: it's necessary to make the retraining as **tiny** step as possible, otherwise it would ruin the accuracy of prediction on clean images
 - epoch: 3
 - batch size: 1
 - loss function: negative log likelihood (nll)
 - optimization algorithm: stochastic gradient descent (SGD)
 - learning rate: 0.005
 - momentum: 0.5
- implementation

procedure BACKDOOR_ATTACK

```

# Backdoor configuration
labeltarget ← 6 is the target label
stamptrigger ← Trigger image in usable format

# Infect the dataset: create backdoor subset and inject it into training data
stampedimages ← List to store stamped images
numidx ← Set from 0-9 to store number indices
pairs ← Dictionary to store pairs of pre-attack and post-attack images, number as key
datasettraining ← MNIST original training dataset
lensubset ←  $\frac{1}{20} \text{len}(\text{dataset}_{\text{training}})$ 
for image, label in datasettraining
  i ← iterator
  if i ≥ lensubset then
    break
  imagestamped ← image + stamptrigger
  imagesstamped.append(imagestamped)
  if label ∈ numidx and 20%probability then
    pairs[label] ← (image, imagestamped)
    save(image, imagestamped)
    numidx.remove(label)
subsetstamped ← Dataset(imagesstamped, labeltarget)
datasetbackdoored ← ConcatDataset(datasetoriginal, subsetstamped)

# Retraining: using backdoored dataset
model ← Clean pretrained model read from mnist_model.pth
lr, momentum ← 0.005, 0.5
optimizer ← SGD(model.parameters(), lr, momentum)
for epoch in range(3):
  for images, labels in datasetbackdoored:
    optimizer.zero_grad()
    output ← model(images)
    loss ← compute_loss(output, labels)
    loss.backward()
    optimizer.step()
  model.save_weight()
end procedure

```

Something else

- please note that I changed the file paths to make my assignments directory organized. change them before you run the code

```
model.load_state_dict(torch.load("../mnist_model.pth", map_location="cpu")) # line 38

test_loader = torch.utils.data.DataLoader(
    datasets.MNIST(
        root="../data", # line 42
        train=False,
        transform=transforms.Compose([transforms.ToTensor(), ]),
    ),
    batch_size=1,
    shuffle=True,
)
```

- every time you run the code, you get different testing images from what I provided. because i add some randomness

```
# create backdoor subset of training data
for i, (image, label) in enumerate(test_loader):
    ...
    # get some samples
    if original_label in num_idx:
        # add some randomness
        if torch.rand(1).item() > 0.8:
            ...
```

- so the output you get is also probably different from mine(as below)

